



FORNAX


How do the RBMs work?

Krzysztof Kolasiński March - 2017

WWW.FORNAX.CO

What is RBM

It is a Boltzmann machine with restriction



What is RBM

RBM's are defined in term of energy function




Fig. 5. The network graph of an RBM with n hidden and m visible units.

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i$$

the random variables (\mathbf{V}, \mathbf{H}) take values $(\mathbf{v}, \mathbf{h}) \in \{0, 1\}^{m+n}$

$\mathbf{V} = (V_1, \dots, V_m)$ - represent the observable data

$\mathbf{H} = (H_1, \dots, H_n)$ - to capture the dependencies between observed variables

\mathbf{W} is a matrix which connects the visible and hidden units

\mathbf{b}, \mathbf{c} are the bias terms associated with visible and hidden units



What is RBM

RBM's are defined in term of energy function




Fig. 5. The network graph of an RBM with n hidden and m visible units.

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i$$

the random variables (\mathbf{V}, \mathbf{H}) take values $(\mathbf{v}, \mathbf{h}) \in \{0, 1\}^{m+n}$

For a given configuration of units \mathbf{V} and \mathbf{H} we can compute energy as:

$$E = -\mathbf{h}^T \mathbf{W} \mathbf{v} - \mathbf{h}^T \mathbf{b} - \mathbf{v}^T \mathbf{c}$$

```
function rbm_energy(v, h, rbm::RBM)
    return -(v' * rbm.w * h + h' * rbm.bh + v' * rbm.bv)
end
```


What is RBM

RBM's are defined in term of energy function




Fig. 5. The network graph of an RBM with n hidden and m visible units.

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i$$

the random variables (\mathbf{V}, \mathbf{H}) take values $(\mathbf{v}, \mathbf{h}) \in \{0, 1\}^{m+n}$

For a given configuration of units \mathbf{V} and \mathbf{H} we can compute energy as:

$$E = -\mathbf{h}^T \mathbf{W} \mathbf{v} - \mathbf{h}^T \mathbf{b} - \mathbf{v}^T \mathbf{c}$$

Each configuration of using has defined probability (Boltzmann distribution)

$$p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})}$$

where $Z = \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$ is a partition function

On partition function

It contains sum of all possible hidden/visible configurations

$$Z = \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

Let us consider the case where $\mathbf{V} = \{v_1, v_2, v_3\}$ and $\mathbf{H} = \{h_1, h_2\}$.

For \mathbf{V} we have 2^3 possible configurations: $\{0,0,0\}, \{0,0,1\}, \{0,1,0\}, \{0,1,1\}, \{1,0,0\}, \{1,0,1\}, \{1,1,0\}, \{1,1,1\}$

For \mathbf{H} we have 2^2 possible configurations: $\{0,0\}, \{0,1\}, \{1,0\}, \{1,1\}$

In real world examples the Z is intractable, for example MNIST images has dimensions of 28x28 pixels. This gives $2^{28 \times 28}$ possible configurations.

On partition function

It contains sum of all possible hidden/visible configurations

$$Z = \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

Let us consider the case where $\mathbf{V} = \{v_1, v_2, v_3\}$ and $\mathbf{H} = \{h_1, h_2\}$.

For \mathbf{V} we have 2^3 possible configurations: $\{0,0,0\}, \{0,0,1\}, \{0,1,0\}, \{0,1,1\}, \{1,0,0\}, \{1,0,1\}, \{1,1,0\}, \{1,1,1\}$

For \mathbf{H} we have 2^2 possible configurations: $\{0,0\}, \{0,1\}, \{1,0\}, \{1,1\}$

In the example above the sum can be computed easily:

$$Z = e^{-E(\{0,0,0\}, \{0,0\})} + e^{-E(\{0,0,0\}, \{0,1\})} + \dots + e^{-E(\{1,1,1\}, \{1,1\})}$$

```
vp = get_permutations(length(rbm.bv))  
hp = get_permutations(length(rbm.bh))  
Z = Z_exact(rbm, hp, vp)
```

```
function Z_exact(rbm::RBM, hp, vp)  
    Z = 0 # partition function  
    for i = 1:length(hp)  
        for j = 1:length(vp)  
            Z += exp(-rbm_energy(vp[j], hp[i], rbm))  
        end  
    end  
    return Z  
end
```

What do we want to maximize?

The RBM: $E = -\mathbf{h}^T \mathbf{W} \mathbf{v} - \mathbf{h}^T \mathbf{b} - \mathbf{v}^T \mathbf{c}$ $Z = \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$

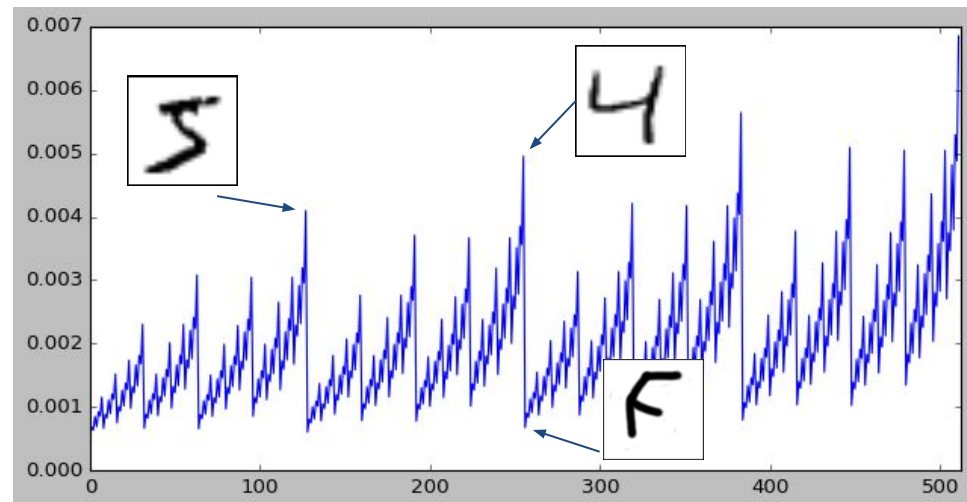
$$p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})}$$

We want to find \mathbf{W} , \mathbf{b} , \mathbf{c} such that the probability of observing \mathbf{V} will be maximal

marginal distribution of \mathbf{V} \rightarrow $p(\mathbf{v}) = \sum_{\mathbf{h}} p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$



Example of the concept: high probabilities for pictures similar as in data set



What do we want to maximize?

The RBM: $E = -\mathbf{h}^T \mathbf{W} \mathbf{v} - \mathbf{h}^T \mathbf{b} - \mathbf{v}^T \mathbf{c}$ $Z = \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$

$$p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})}$$

We want to find \mathbf{W} , \mathbf{b} , \mathbf{c} such that the probability of observing \mathbf{V} will be maximal

$$p(\mathbf{v}) = \sum_{\mathbf{h}} p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

We look for optimal $(\mathbf{W}, \mathbf{b}, \mathbf{c}) = \boldsymbol{\theta}$ using standard **gradient ascent** on log-likelihood

$$\ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v}) = \ln p(\mathbf{v} | \boldsymbol{\theta}) = \ln \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} = \ln \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} - \ln \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} + \underbrace{\eta \frac{\partial}{\partial \boldsymbol{\theta}^{(t)}} \left(\ln \mathcal{L}(\boldsymbol{\theta}^{(t)} | S) \right)}_{= \Delta \boldsymbol{\theta}^{(t)}} - \lambda \boldsymbol{\theta}^{(t)} + \nu \Delta \boldsymbol{\theta}^{(t-1)} \quad \text{Regularization + momentum}$$

Expression for gradients

We are looking for expression of gradient of L with respect to model parameters (**theta**)

$$\ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v}) = \ln p(\mathbf{v} | \boldsymbol{\theta}) = \ln \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} = \ln \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} - \ln \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

It's a simple math:

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial \boldsymbol{\theta}} = \frac{\partial}{\partial \boldsymbol{\theta}} \left(\ln \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} \right) - \frac{\partial}{\partial \boldsymbol{\theta}} \left(\ln \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} \right)$$

$$= - \sum_{\mathbf{h}} p(\mathbf{h} | \mathbf{v}) \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial \boldsymbol{\theta}} + \sum_{\mathbf{v}, \mathbf{h}} p(\mathbf{v}, \mathbf{h}) \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial \boldsymbol{\theta}}$$

$$p(\mathbf{h} | \mathbf{v}) = \frac{p(\mathbf{v}, \mathbf{h})}{p(\mathbf{v})} = \frac{\frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})}}{\frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}} = \frac{e^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}}$$

$$p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})}$$

This is still intractable

Finding expression for $p(\mathbf{h} | \mathbf{v})$

Due to the structure of RBMs the conditional probability $p(\mathbf{h} | \mathbf{v})$ can be found analytically

$$- \sum_{\mathbf{h}} \boxed{p(\mathbf{h} | \mathbf{v})} \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial \boldsymbol{\theta}}$$

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i$$

Let's group E by h_i

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i^n h_i \left(\sum_j^m \omega_{ij} v_j + c_i \right) - \sum_j^m b_j v_j$$

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i^n h_i \Delta_i - E_v$$

Let's write it explicitly

$$e^{-E(\mathbf{v}, \mathbf{h})} = e^{\sum_i^n h_i \Delta_i + E_v} = e^{\sum_i^n h_i \Delta_i} e^{E_v} = e^{E_v} \prod_i^n e^{h_i \Delta_i}$$

$$\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} = \sum_{\mathbf{h}} e^{E_v} \prod_i^n e^{h_i \Delta_i} = e^{E_v} \sum_{\mathbf{h}} \prod_i^n e^{h_i \Delta_i}$$

$$e^{E_v} \sum_{\mathbf{h}} \prod_i^n e^{h_i \Delta_i} = e^{E_v} \sum_{h_1} \sum_{h_2} \dots \sum_{h_n} e^{h_1 \Delta_1} e^{h_2 \Delta_2} \dots e^{h_n \Delta_n}$$

$$= e^{E_v} \sum_{h_1} e^{h_1 \Delta_1} \sum_{h_2} e^{h_2 \Delta_2} \dots \sum_{h_n} e^{h_n \Delta_n} =$$

Finding expression for $p(\mathbf{h} | \mathbf{v})$

Due to the structure of RBMs the conditional probability $p(\mathbf{h} | \mathbf{v})$ can be found analytically

$$- \sum_{\mathbf{h}} \boxed{p(\mathbf{h} | \mathbf{v})} \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial \boldsymbol{\theta}}$$

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i$$

Let's group E by h_i

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i h_i \left(\sum_j \omega_{ij} v_j + c_i \right) - \sum_j b_j v_j$$

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i h_i \Delta_i - E_v$$

$$p(\mathbf{h} | \mathbf{v}) = \frac{p(\mathbf{v}, \mathbf{h})}{p(\mathbf{v})} = \frac{\frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})}}{\frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}} = \frac{e^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}}$$

Let's write it explicitly

$$e^{-E(\mathbf{v}, \mathbf{h})} = e^{E_v} \sum_{h_1} e^{h_1 \Delta_1} \sum_{h_2} e^{h_2 \Delta_2} \dots \sum_{h_n} e^{h_n \Delta_n} =$$

h_i is a binary variable: $h_i = \{0, 1\}$

$$\sum_{h_i} e^{h_i \Delta_i} = 1 + e^{\Delta_i}$$

This gives

$$\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} = e^{E_v} (1 + e^{\Delta_1}) \dots (1 + e^{\Delta_n}) = e^{E_v} \prod_i^n (1 + e^{\Delta_i})$$

Finding expression for $p(\mathbf{h} | \mathbf{v})$

Due to the structure of RBMs the conditional probability $p(\mathbf{h} | \mathbf{v})$ can be found analytically

$$- \sum_{\mathbf{h}} \boxed{p(\mathbf{h} | \mathbf{v})} \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial \boldsymbol{\theta}}$$

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i$$

Let's group E by h_i

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i h_i \left(\sum_j w_{ij} v_j + c_i \right) - \sum_j b_j v_j$$

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i h_i \Delta_i - E_v$$

$$p(\mathbf{h} | \mathbf{v}) = \frac{p(\mathbf{v}, \mathbf{h})}{p(\mathbf{v})} = \frac{\frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})}}{\frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}} = \frac{e^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}}$$

Finally we have

$$\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} = e^{E_v} \prod_i (1 + e^{\Delta_i}) \quad e^{-E(\mathbf{v}, \mathbf{h})} = e^{E_v} \prod_i e^{h_i \Delta_i}$$

From this we can compute conditional probability

$$p(\mathbf{h} | \mathbf{v}) = \frac{e^{E_v} \prod_i e^{h_i \Delta_i}}{e^{E_v} \prod_i (1 + e^{\Delta_i})} = \prod_i \boxed{\frac{e^{h_i \Delta_i}}{1 + e^{\Delta_i}}}$$

$$p(\mathbf{h} | \mathbf{v}) = \prod_i \boxed{p(h_i | \mathbf{v})}$$

Probability of activation of i-th hidden unit given visible units \mathbf{v}

$$p(h_i = 1 | \mathbf{v}) = \frac{e^{\Delta_i}}{1 + e^{\Delta_i}} = \frac{1}{1 + e^{-\Delta_i}} = \text{sig}(\Delta_i)$$

Finding expression for $p(\mathbf{h} | \mathbf{v})$

Due to the structure of RBMs the conditional probability $p(\mathbf{h} | \mathbf{v})$ can be found analytically

$$- \sum_{\mathbf{h}} \boxed{p(\mathbf{h} | \mathbf{v})} \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial \boldsymbol{\theta}}$$

$$p(\mathbf{h} | \mathbf{v}) = \prod_i^n p(h_i | \mathbf{v})$$

$$\begin{aligned} p(h_i = 1 | \mathbf{v}) &= \text{sig}(+\Delta_i) \\ p(h_i = 0 | \mathbf{v}) &= \text{sig}(-\Delta_i) \end{aligned}$$

$$\Delta_i = \sum_j^m \omega_{ij} v_j + c_i$$

Conclusions:

- *RBM works as stochastic neuron with sigmoid activation function*
- *Each hidden unit is independent (restriction in hidden-hidden connections)*
- *Similarly for visible units*

$$p(H_i = 1 | \mathbf{v}) = \text{sig} \left(\sum_{j=1}^m w_{ij} v_j + c_i \right)$$

$$p(V_j = 1 | \mathbf{h}) = \text{sig} \left(\sum_{i=1}^n w_{ij} h_i + b_j \right)$$




Fig. 5. The network graph of an RBM with n hidden and m visible units.

The gradient of L

The gradient can be computed using similar factorization method

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i$$

After few steps of calculations we finally obtain

$$\sum_{\mathbf{h}} p(\mathbf{h} | \mathbf{v}) \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial w_{ij}} = \text{sig} \left(\sum_{j=1}^m w_{ij} v_j + c_i \right) v_j$$

$$p(h_i = 1 | \mathbf{v}) = \text{sig} (+\Delta_i)$$

$$p(h_i = 0 | \mathbf{v}) = \text{sig} (-\Delta_i)$$

$$\Delta_i = \sum_j^m \omega_{ij} v_j + c_i$$

The final expression for the gradient of log-likelihood is given:

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial w_{ij}} = - \sum_{\mathbf{h}} p(\mathbf{h} | \mathbf{v}) \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial w_{ij}} + \sum_{\mathbf{v}, \mathbf{h}} p(\mathbf{v}, \mathbf{h}) \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial w_{ij}}$$

$$= \sum_{\mathbf{h}} p(\mathbf{h} | \mathbf{v}) h_i v_j - \sum_{\mathbf{v}} p(\mathbf{v}) \sum_{\mathbf{h}} p(\mathbf{h} | \mathbf{v}) h_i v_j = p(H_i = 1 | \mathbf{v}) v_j - \sum_{\mathbf{v}} p(\mathbf{v}) p(H_i = 1 | \mathbf{v}) v_j$$

This is still intractable

The gradient of L (per one example)

Final expressions:

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial w_{ij}} = p(H_i = 1 | \mathbf{v}) v_j - \sum_{\mathbf{v}} p(\mathbf{v}) p(H_i = 1 | \mathbf{v}) v_j$$

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial b_j} = v_j - \sum_{\mathbf{v}} p(\mathbf{v}) v_j$$

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial c_i} = p(H_i = 1 | \mathbf{v}) - \sum_{\mathbf{v}} p(\mathbf{v}) p(H_i = 1 | \mathbf{v})$$

probability of firing rate of the neuron

$$p(H_i = 1 | \mathbf{v}) = \text{sig} \left(\sum_{j=1}^m w_{ij} v_j + c_i \right)$$

marginal probability

$$p(\mathbf{v}) = \sum_{\mathbf{h}} p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

probability of \mathbf{v} generated by model

Minibatch - what is at equilibrium?

For **mini-batch** algorithm one computes

$$\frac{1}{\ell} \sum_{\mathbf{v} \in S} \frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial w_{ij}}$$

Equilibrium condition: *consider $\ell = \text{infinity}$, and gradient zero*

$$\frac{1}{\ell} \sum_{\mathbf{v} \in S} \frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial w_{ij}} = \frac{1}{\ell} \sum_{\mathbf{v} \in S} p(h_i = 1 | \mathbf{v}) v_j - \sum_{\mathbf{v}'} p(\mathbf{v}') p(h_i = 1 | \mathbf{v}') v'_j = \sum_{\mathbf{v}} q(\mathbf{v}) p(h_i = 1 | \mathbf{v}) v_j - \sum_{\mathbf{v}'} p(\mathbf{v}') p(h_i = 1 | \mathbf{v}') v'_j = 0$$

$$\sum_{\mathbf{v}} p(h_i = 1 | \mathbf{v}) v_j (q(\mathbf{v}) - p(\mathbf{v})) = 0 \implies q(\mathbf{v}) = p(\mathbf{v})$$

for large mini-batch model and data distributions should be the same

Approximating gradients by sampling

The gradient with respect to \mathbf{b} vector

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial b_j} = v_j - \sum_{\mathbf{v}} p(\mathbf{v}) v_j$$

expectation value of \mathbf{v}
over the model probability

We do not know nice and
compact form of $p(\mathbf{v})$:(

Let's us recall simple example of sampling from gaussian distribution $\rho(\mathbf{x})$

$$\mathbb{E}_{\rho(x)} [f(x)] = \int_{-\infty}^{+\infty} \rho(x) f(x) dx$$

This integral can be approximated by average of sampled x from ρ distribution

$$\mathbb{E}_{\rho(x)} [f(x)] \approx \frac{1}{N} \sum_{x_i \sim \rho(x)} f(x_i)$$

x is sampled from $\rho(x)$

A toy example of sampling approach

Consider following example:

$$\rho(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}} \quad f(x) = x^2$$

$$\mathbb{E}_{\rho(x)} [f(x)] = \int_{-\infty}^{+\infty} \rho(x) f(x) dx = 1$$

Approximation by sampling x from rho

$$\mathbb{E}_{\rho(x)} [f(x)] \approx \frac{1}{N} \sum_{x_i \sim \rho(x)} f(x_i)$$

```
for n = 1:6
    E = mean((randn(10^n).^2))
    println("${10^n}\t$tE")
end
```

N	E(x^2)
10	0.351830
100	1.046075
1000	0.939235
10000	1.027309
100000	1.001297
1000000	0.999072

In the example above we have sampled x from rho, however in real world examples we don't know exact rho or we cannot easily sample from rho
;(
(

Approximating gradients by sampling

The gradient with respect to \mathbf{b} vector

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial b_j} = v_j - \sum_{\mathbf{v}} p(\mathbf{v}) v_j$$

Let's come back to the RBMs case:

$$\sum_{\mathbf{v}} p(\mathbf{v}) f(\mathbf{v}) = \mathbb{E}_{p(\mathbf{v})} [f(\mathbf{v})] \approx \frac{1}{N} \sum_{\mathbf{v} \sim p(\mathbf{v})} f(\mathbf{v})$$

N is a number of samples generated by $p(\mathbf{v})$

We want to approximate expectation value with sampling.

The only question which left is:

How to sample from complex distribution $p(\mathbf{v})$?

$$\begin{aligned} \frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial w_{ij}} &= p(H_i = 1 | \mathbf{v}) v_j - \sum_{\mathbf{v}} p(\mathbf{v}) p(H_i = 1 | \mathbf{v}) v_j \\ \frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial b_j} &= v_j - \sum_{\mathbf{v}} p(\mathbf{v}) v_j \\ \frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial c_i} &= p(H_i = 1 | \mathbf{v}) - \sum_{\mathbf{v}} p(\mathbf{v}) p(H_i = 1 | \mathbf{v}) \end{aligned}$$

Markov chain Monte Carlo method (MCMCM)

We want to approximate:

$$\sum_{\mathbf{v}} p(\mathbf{v}) f(\mathbf{v}) = \mathbb{E}_{p(\mathbf{v})} [f(\mathbf{v})] \approx \frac{1}{N} \sum_{\mathbf{v} \sim p(\mathbf{v})} f(\mathbf{v})$$

This can be done via sampling method

Some definitions:

- **Markov chain** is a time discrete stochastic process, where the next state of the system depends only on current state of the system

$$p_{ij}^{(k)} = \Pr \left(X^{(k+1)} = j \mid X^{(k)} = i, X^{(k-1)} = i_{k-1}, \dots, X^{(0)} = i_0 \right) = \Pr \left(X^{(k+1)} = j \mid X^{(k)} = i \right)$$

memoryless process

next state current state

- Many more: irreducible, aperiodic, detailed balance condition.... to define when we can use MCMC methods (stationary solution)

Markov chain can be generated with Metropolis sampling however this has some drawbacks:
states can be rejected :(

Solution? Gibbs sampling method - no rejections!

Gibbs sampling method for RBMs

We want to approximate:

$$\sum_{\mathbf{v}} p(\mathbf{v}) f(\mathbf{v}) = \mathbb{E}_{p(\mathbf{v})} [f(\mathbf{v})] \approx \frac{1}{N} \sum_{\mathbf{v} \sim p(\mathbf{v})} f(\mathbf{v})$$

This can be done via Gibbs sampling method:

- if the target distribution is $p(\mathbf{x})$ we must have analytical expressions for conditional probabilities: $p(x_i | \mathbf{x}_{-i})$ for $j < i$


$$p(H_i = 1 | \mathbf{v}) = \text{sig} \left(\sum_{j=1}^m w_{ij} v_j + c_i \right)$$

Done!

$$p(V_j = 1 | \mathbf{h}) = \text{sig} \left(\sum_{i=1}^n w_{ij} h_i + b_j \right)$$

- we must be able to sample from those probabilities e.g if $p(h=1|\mathbf{v}) = 0.9 > \text{rand}() \Rightarrow h=1$ else $h=0$ Done!

The algorithm:



$$h^{(n+1)} \sim \text{sigm}(W'v^{(n)} + c)$$
$$v^{(n+1)} \sim \text{sigm}(Wh^{(n+1)} + b),$$

after some steps \mathbf{v}^k is a sample from $p(\mathbf{v})$ distribution.

Gibbs sampling method for RBMs

We want to approximate:

$$\sum_{\mathbf{v}} p(\mathbf{v}) f(\mathbf{v}) = \mathbb{E}_{p(\mathbf{v})} [f(\mathbf{v})] \approx \frac{1}{N} \sum_{\mathbf{v} \sim p(\mathbf{v})} f(\mathbf{v})$$

Let's back to RBM we can compute $p(\mathbf{v})$ using exact approach

```
rbm = RBM(9, 9)
vp = get_permutations(length(rbm.bv))
hp = get_permutations(length(rbm.bh))
Z = Z_exact(rbm, hp, vp) | 1.13e6...
```

```
probs_pv = zeros(length(vp))
for s = 1:length(vp)
    probs_pv[s] = p_v_exact(vp[s], Z, hp, rbm)
end | ✓
```

Let's sample $p(\mathbf{v})$ with Gibbs method


```
vk = copy(samples[1,:]) # initialize with data sample
gibbs_pv = zeros(nvp)
k = 100000
for i = 1:k
    hk = sample_bernoulli(p_h_cond_v(vk, rbm))
    vk = sample_bernoulli(p_v_cond_h(hk, rbm))
    idx = parse{Int64, join(vk), 2} + 1 # get idx
    gibbs_pv[idx] += 1.0/k
end | ✓
```

```
function sample_bernoulli(probs)
    return 1(rand(length(probs)) .< probs)
end
```

p_v_exact:


$$p(\mathbf{v}) = \sum_{\mathbf{h}} p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

Gibbs sampling



$$p(H_i = 1 | \mathbf{v}) = \text{sig} \left(\sum_{j=1}^m w_{ij} v_j + c_i \right)$$

$$p(V_j = 1 | \mathbf{h}) = \text{sig} \left(\sum_{i=1}^n w_{ij} h_i + b_j \right)$$




Gibbs sampling method for RBMs

We want to approximate:

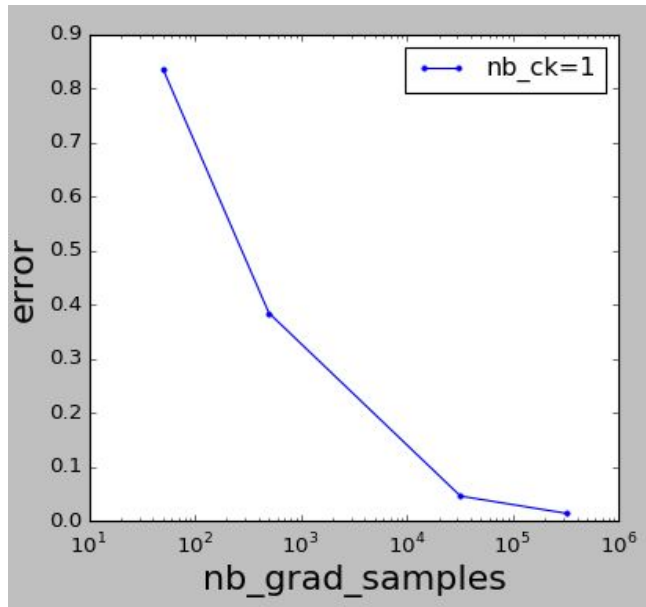
$$\sum_{\mathbf{v}} p(\mathbf{v}) f(\mathbf{v}) = \mathbb{E}_{p(\mathbf{v})} [f(\mathbf{v})] \approx \frac{1}{N} \sum_{\mathbf{v} \sim p(\mathbf{v})} f(\mathbf{v})$$

Gibbs sampling allows us to sample $\mathbf{v} \sim p(\mathbf{v})$

Gibbs sampling



How many samples k should we compute to estimate gradients properly?




Relative error of gradient computation using Gibbs sampling and exact

What people do?

They use $k = 1$

Contrastive divergence



The idea of k -step contrastive divergence learning (CD- k) is quite simple: instead of approximating the second term in the log-likelihood gradient by a sample from the RBM-distribution (which would require running a Markov chain until the stationary distribution is reached), a Gibbs chain is run for only k steps (and usually $k = 1$). The Gibbs chain is initialized with a training example $\mathbf{v}^{(0)}$ of the training set and yields the sample $\mathbf{v}^{(k)}$ after k steps. Each step t consists of sampling $\mathbf{h}^{(t)}$ from

Contrastive divergence

Algorithm 1: k -step contrastive divergence

Input: RBM $(V_1, \dots, V_m, H_1, \dots, H_n)$, training batch S

Output: gradient approximation Δw_{ij} , Δb_j and Δc_i for $i = 1, \dots, n$, $j = 1, \dots, m$

```
1 init  $\Delta w_{ij} = \Delta b_j = \Delta c_i = 0$  for  $i = 1, \dots, n$ ,  $j = 1, \dots, m$ 
2 forall the  $v \in S$  do
3    $v^{(0)} \leftarrow v$ 
4   for  $t = 0, \dots, k - 1$  do
5     for  $i = 1, \dots, n$  do sample  $h_i^{(t)} \sim p(h_i | v^{(t)})$ 
6     ;
7     for  $j = 1, \dots, m$  do sample  $v_j^{(t+1)} \sim p(v_j | h^{(t)})$ 
8     ;
9   for  $i = 1, \dots, n$ ,  $j = 1, \dots, m$  do
10     $\Delta w_{ij} \leftarrow \Delta w_{ij} + p(H_i = 1 | v^{(0)}) \cdot v_j^{(0)} - p(H_i = 1 | v^{(k)}) \cdot v_j^{(k)}$ 
11  for  $j = 1, \dots, m$  do
12     $\Delta b_j \leftarrow \Delta b_j + v_j^{(0)} - v_j^{(k)}$ 
13  for  $i = 1, \dots, n$  do
14     $\Delta c_i \leftarrow \Delta c_i + p(H_i = 1 | v^{(0)}) - p(H_i = 1 | v^{(k)})$ 
```

This seems to work quite well but for small values of k one has to be aware of bad approximation of gradients - gradient is biased (see more in refs).

$$\frac{\partial \ln \mathcal{L}(\theta | v)}{\partial b_j} = v_j - \sum_v p(v) v_j$$

Persistent contrastive divergence (PCD)

Algorithm 1: k -step contrastive divergence

Input: RBM $(V_1, \dots, V_m, H_1, \dots, H_n)$, training batch S

Output: gradient approximation Δw_{ij} , Δb_j and Δc_i for $i = 1, \dots, n$, $j = 1, \dots, m$

```
1 init  $\Delta w_{ij} = \Delta b_j = \Delta c_i = 0$  for  $i = 1, \dots, n$ ,  $j = 1, \dots, m$ 
2 forall the  $v \in S$  do
3    $v^{(0)} \leftarrow v$ 
4   for  $t = 0, \dots, k - 1$  do
5     for  $i = 1, \dots, n$  do sample  $h_i^{(t)} \sim p(h_i | v^{(t)})$ 
6     ;
7     for  $j = 1, \dots, m$  do sample  $v_j^{(t+1)} \sim p(v_j | h^{(t)})$ 
8     ;
9   for  $i = 1, \dots, n$ ,  $j = 1, \dots, m$  do
10     $\Delta w_{ij} \leftarrow \Delta w_{ij} + p(H_i = 1 | v^{(0)}) \cdot v_j^{(0)} - p(H_i = 1 | v^{(k)}) \cdot v_j^{(k)}$ 
11  for  $j = 1, \dots, m$  do
12     $\Delta b_j \leftarrow \Delta b_j + v_j^{(0)} - v_j^{(k)}$ 
13  for  $i = 1, \dots, n$  do
14     $\Delta c_i \leftarrow \Delta c_i + p(H_i = 1 | v^{(0)}) - p(H_i = 1 | v^{(k)})$ 
```

don't re-initialize $v^{(0)}$ for each example
in mini-batch

Dependence on k

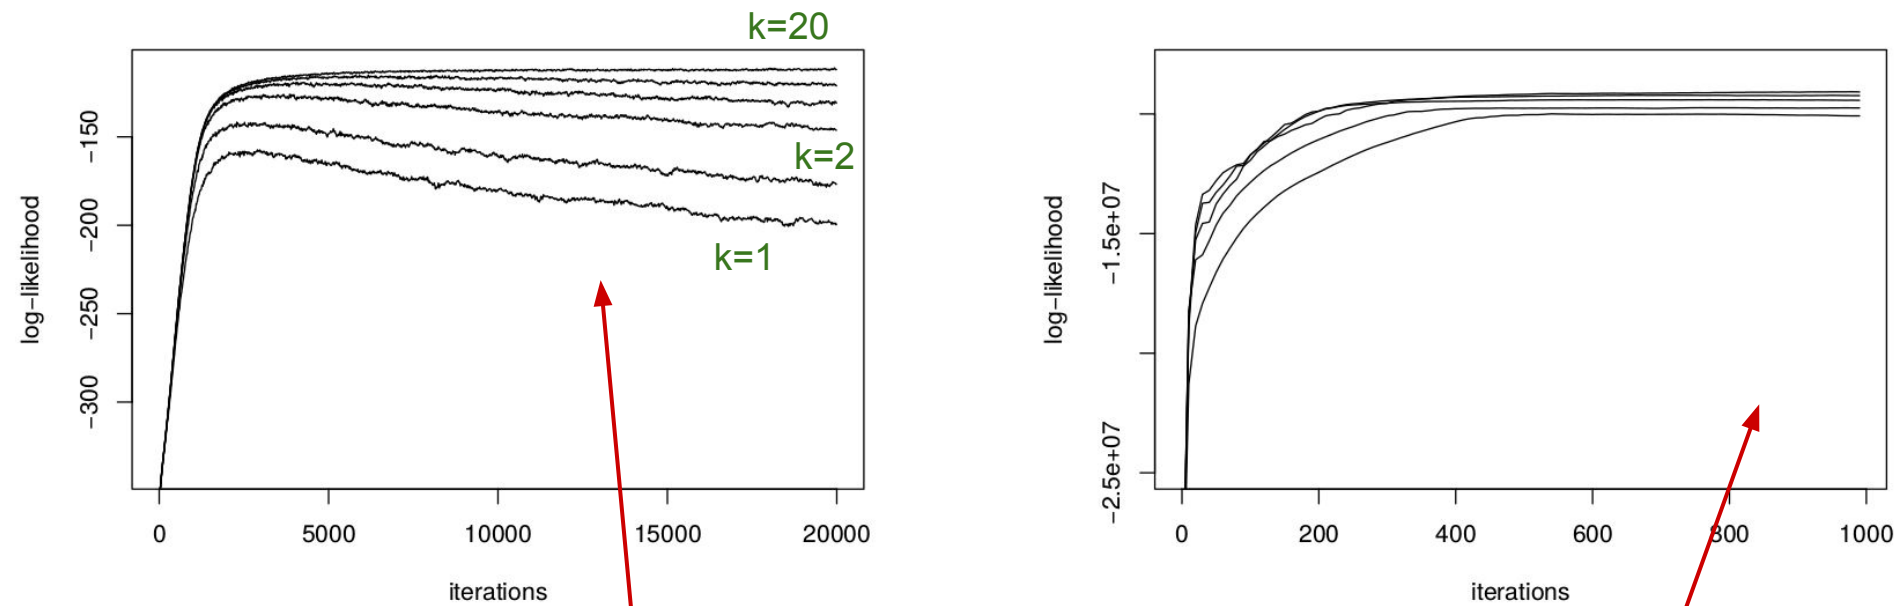


Fig. 9. Evolution of the log-likelihood during training of RBMs with CD- k where different values for k were used. The left plot shows the results for BAS (from bottom to top $k = 1, 2, 5, 10, 20, 100$) and the right plot for MNIST (from bottom to top $k = 1, 2, 5, 10, 20$). The values are medians over 25 runs.




Fig. 6. Top: Patterns from the BAS data set. Bottom: Images from the MNIST data set.

Summary




Fig. 5. The network graph of an RBM with n hidden and m visible units.

RBM is an energy based model

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i$$

We want to maximize log likelihood of $p(\mathbf{v})$

$$p(\mathbf{v}) = \sum_{\mathbf{h}} p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

Using gradient ascent and CD approximation method


$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial w_{ij}} = p(H_i = 1 | \mathbf{v}) v_j - \sum_{\mathbf{v}} p(\mathbf{v}) p(H_i = 1 | \mathbf{v}) v_j$$

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial b_j} = v_j - \sum_{\mathbf{v}} p(\mathbf{v}) v_j$$

$$\frac{\partial \ln \mathcal{L}(\boldsymbol{\theta} | \mathbf{v})}{\partial c_i} = p(H_i = 1 | \mathbf{v}) - \sum_{\mathbf{v}} p(\mathbf{v}) p(H_i = 1 | \mathbf{v})$$

Applications of RBMs

- RBMs can be used to pretrain weights for supervised learning




- RBMs are generative model, one can solve in-painting problem with them

$$p(V_{o+1}, \dots, V_m \mid V_1 = v_1, \dots, V_o = v_o)$$

keep selected units fixed and sample rest using Gibbs sampling

- Train Deep Belief networks
- RBMs were used for recommendations :)



Notes: Why don't we work with Energy?

- RBM is an energy based model, why not minimize energy $E(\mathbf{v})$ instead of $p(\mathbf{v})$


$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i \longrightarrow p(\mathbf{v}) = \sum_{\mathbf{h}} p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

Smaller energy => higher probability

$$E(\mathbf{v}) = \frac{1}{N_h} \sum_{\mathbf{h}} E(\mathbf{v}, \mathbf{h})$$

- This is not true in general for $E(\mathbf{v})$
- Energy is not bounded, probability must be normalized

Results for *converged* 3x3 BAS dataset



Conclusions

Don't use RBMs if you don't have to

References

- <https://theclevermachine.wordpress.com/page/3/> -mcmc
- <https://theclevermachine.wordpress.com/page/2/> gibbs sampling
- <http://deeplearning.net/tutorial/rbm.html> - theano tutorial on RBMs (full implementation)
- [Training RBMs: An introduction](#) - the best description on how do the RBMs work



FORNAX

WWW.FORNAX.CO